
Entry Without Displacement: Open-Weight Models and the Market for Inference

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Open-weight models are widely expected to constrain the pricing power of propri-
2 etary model providers, but the mechanism has not been measured. We assemble
3 a market-wide panel of the models served by a large inference router, covering
4 posted prices, capability benchmarks, weight-release status and revealed usage
5 for 457 models and extended back to 2023 with archived catalogues. Competition
6 works through releases rather than repricing. 54% of models are never repriced
7 at all, and a posted price changes in only 12.6% of model-months, even though
8 changing it costs a provider nothing. When prices do change, volume responds
9 weakly. The own-price elasticity is -0.53 (s.e. 0.10), identified from 708 revisions
10 with model and date fixed effects. And entry does not displace incumbents.
11 Over four weeks, models released inside the window took 28% of tokens while
12 incumbent volume grew 18%, and we can rule out differential losses above 0.13
13 log points for the incumbents closest to an entrant in capability space. Removing
14 open-weight models from the choice set in an estimated nested-logit demand system
15 raises the average price paid per token by 156% and costs users about \$2.4m
16 per day on this platform, while raising provider concentration. Open weights discipline
17 the price at which new demand is served, not the prices incumbents post.

18 1 Introduction

19 Whether open-weight models constrain the market power of proprietary providers is now a policy
20 question as much as an economic one. Reports commissioned by governments name it as an open
21 evidence gap [Bengio et al., 2026], and the theoretical case cuts both ways. Scale economies in training
22 push toward concentration [Kaplan et al., 2020, Korinek and Vipra, 2025], while downloadable
23 weights supply a competitive fringe that any host can serve, the mechanism open-source software
24 has long been argued to supply in adjacent markets [Lerner and Tirole, 2002, Nagle, 2019]. Demirer
25 et al. [2025] document a collapse in the price of a unit of measured intelligence and a proliferation
26 of suppliers, and Nagle and Yue [2025] show that much proprietary usage sits on models dominated
27 on price and measured capability by an open alternative.

28 Both results describe the allocation of demand across models at a point in time. Neither identifies
29 the channel through which open weights would constrain proprietary pricing, and the two candidates
30 have opposite implications for policy. If open weights force incumbents to cut posted prices, the
31 constraint is visible in the price series of established models. If instead open weights capture new
32 demand while incumbents hold their prices, the constraint operates on the terms at which the market
33 grows, and any measurement built on incumbent price responses will find nothing.

34 We assemble a panel that separates the two, combining the current cross-section of a large inference
35 router (457 models with observed usage, 175 of them open-weight), a daily usage panel covering all
36 of them for one month, task-level shares for 29 uses, per-provider hosting terms, and archived price
37 catalogues and leaderboards reaching back to 2023. The panel used for price responses contains
38 19,803 model-date observations on 634 models. Three results follow, and they point the same way.
39 First, the posted price of a given model version behaves like a fixed characteristic of the product. 54%

of models are never repriced, and the monthly frequency of a price change is 12.6%, comparable to consumer goods [Nakamura and Steinsson, 2008, Klenow and Malin, 2011] in a setting where the menu cost that explains rigidity there does not exist. Second, when prices do change, volume moves little. Exploiting 708 revisions with model and date fixed effects, we estimate an own-price elasticity of -0.53 (s.e. 0.10), below the value near one reported by Demiret et al. [2025] and far below what homogeneous-goods reasoning would imply. Third, entry is absorbed by growth rather than by displacement. Over the four weeks of the daily panel the market grew 65%, models released inside the window took 28% of tokens, and incumbent volume still rose 18%. In an incumbent-by-entrant event study we find no differential decline for the incumbents technologically closest to an entrant.

We then quantify the constraint where it does operate. We estimate a nested-logit demand system over the 172 models with prices and capability scores, invert it to recover mean utilities that reproduce observed shares, and delete open-weight models from the choice set. Users would pay 156% more per token, and the surplus loss is worth about \$2.4m per day on this platform alone. Across 29 task-level markets the implied price increase rises almost one for one with the open-weight share of task spending. Concentration moves against the usual presumption, since removing open weights raises the provider Herfindahl index from 897 to 1048.

Our contribution is to identify the margin on which open-weight competition operates and to price it. The measurement of price rigidity and of the within-model demand response is, to our knowledge, new to this market, and it explains why the reallocation gains documented by Nagle and Yue [2025] persist. With inelastic short-run demand a provider has little to gain from cutting the price of a model already in the field, and users have little reason to move quickly. We release the collection code, which reconstructs the entire panel from public endpoints.

2 Data

2.1 Sources

We use the public endpoints of OpenRouter, a router that serves requests to hundreds of models from many providers through one interface. Four supply the raw material. A catalogue gives posted prices, context length, release timestamp and the Hugging Face repository for models whose weights are published. A leaderboard gives tokens, requests and tool calls by model, a benchmark table gives third-party capability scores, and a task table gives each model’s share of spending in 29 usage categories. Each model’s own page embeds a daily usage series for the trailing month, which we collect for every model with observed usage, and per-provider hosting terms come from the endpoint listing. Historical coverage comes from Internet Archive captures of the catalogue and of the leaderboard, the latter embedding the leaderboard payload in the page’s streamed render data. Appendix A documents each step.

2.2 Openness

We classify model j as open-weight, $o_j = 1$, if the catalogue records a public weight repository for it. The classification is model-specific rather than provider-specific, which matters here because several providers, among them Google, OpenAI and Alibaba, ship both open-weight and proprietary models, and a provider-level rule would misclassify them. The definition follows the one used in the policy literature, where the salient feature is that weights can be downloaded and served by a third party [Kapoor et al., 2024, NTIA, 2024].

2.3 Capability

Capability enters as a vector $q_j \in \mathbb{R}^M$ of third-party scores over M dimensions covering reasoning, coding and agentic use, supplemented by pairwise-comparison scores from a design-oriented arena in the spirit of Chiang et al. [2024]. These measures are highly collinear. Writing $\bar{q}_m$ and ς_m for the mean and standard deviation of dimension m across models, and Λ for the covariance matrix of the standardised scores, the leading eigenvector w of Λ defines the capability index

$$Q_j = w' \tilde{q}_j, \quad \tilde{q}_{jm} = (q_{jm} - \bar{q}_m) / \varsigma_m, \quad (1)$$

which absorbs 98% of the variance of the three benchmark dimensions and 91% when the two arena blocks are added. Where a dimension is missing we impute it by ridge regression on the dimensions a model does have, recovering held-out values with R^2 above 0.96 and giving 180 scored models covering 91% of tokens.

Where a pairwise notion of substitutability is needed we use the whitened distance between capability vectors and the similarity it induces,

$$d_{jk} = \sqrt{(\tilde{q}_j - \tilde{q}_k)' \Lambda^{-1} (\tilde{q}_j - \tilde{q}_k)}, \quad S_{jk} = \exp(-\kappa d_{jk}), \quad (2)$$

with Λ ridge-shrunk and κ normalised so that the median pair maps to e^{-1} . Section 5 reports results under Euclidean and cosine alternatives.

2.4 Prices and quantities

Providers post separate input and output prices, p_{jt}^{in} and p_{jt}^{out} , and models differ sharply in how many output tokens they emit per input token. We therefore work with the blended price

$$p_{jt} = \omega_{jt} p_{jt}^{\text{in}} + (1 - \omega_{jt}) p_{jt}^{\text{out}}, \quad \omega_{jt} = \frac{I_{jt}}{I_{jt} + O_{jt}}, \quad (3)$$

where I_{jt} and O_{jt} are input and output tokens. For the panel regressions we replace ω_{jt} by its sample mean, so the index carries no feedback from a model's own usage composition. Quantities are tokens, $y_{jt} = I_{jt} + O_{jt}$, and spending is $e_{jt} = (p_{jt}^{\text{in}} I_{jt} + p_{jt}^{\text{out}} O_{jt}) / 10^6$.

2.5 Sample and market definition

Table 1 summarises the estimation sample. The platform is one segment of the market for inference, excluding traffic sent directly to a provider's own API and consumer subscription products, and its users are developers who have chosen a router. That selects for price sensitivity and willingness to switch, so the segment we observe is the one in which open-weight substitution is easiest.

Table 1: The estimation sample, measured on the last complete day of the daily panel. Prices are dollars per million tokens for the blended mix of (3). Capability is the index (1). Shares are within the estimation sample, which covers 91% of platform tokens.

	Models	Median price	Median capability	Median context (thousands)	Token share	Spend share
Open-weight	86	0.29	30.5	256	0.73	0.32
Proprietary	86	1.50	36.6	400	0.27	0.68
All scored	172	0.59	33.4	262	1.00	1.00

3 Market structure

3.1 Prices and capability

Figure 1 places every model with usage on the price-capability plane. At any capability level below the frontier, prices differ by one to two orders of magnitude, so posted price is far from a sufficient statistic for what users buy, and the open-weight side is systematically cheaper. Matching each proprietary model to the open-weight models within two points of it on the capability index, the median proprietary model costs 3.3 times the median open alternative, and 1.7 times weighted by volume. Averaged over usage a token from an open-weight model costs \$0.35 per million against \$1.95. These gaps are smaller than the ninety percent difference reported for 2025 data [Demirer et al., 2025, Nagle and Yue, 2025], as one would expect if proprietary providers have since released cheaper tiers.

The right panel tracks the best capability obtainable from each side over time. The open-weight frontier reaches a given capability level a median of 105 days after the proprietary frontier first reaches it. The gap is a lag rather than a level. At the very top of the capability distribution proprietary models are unmatched at every date, but the set of tasks for which that matters shrinks as the open frontier advances.




Figure 1: Left, posted blended price against the capability index, marker area proportional to the square root of token volume, step functions giving the cheapest price available at or above each capability level. Right, the best capability available from each side by release date, shading the open-weight lag.

3.2 Volume shares and revenue shares

Open-weight models account for 68% of tokens but only 30% of spending, and the wedge is visible provider by provider. Figure 2 shows the two distributions side by side. DeepSeek serves 24% of tokens and collects 6% of spending, while Anthropic serves 7% of tokens and collects 45%. The provider Herfindahl index is 1183 computed on tokens and 2439 computed on spending. By volume the market is moderately concentrated. By revenue it is concentrated. Volume shares describe who runs the world’s inference, revenue shares who captures its value.




Figure 2: Left, open-weight share of tokens by quarter, computed from the archived leaderboards. Right, share of platform tokens and of spending for the ten largest providers by volume.

The open-weight share of volume is also far from stable. In the archived leaderboards, open-weight models held 86% of tokens in late 2023, fell to 29% by mid-2025, and recovered to 48% by mid-2026. The trough coincides with the sample used by Nagle and Yue [2025], whose finding that proprietary models dominate usage describes a market that has since moved.

3.3 Supply of open weights

Openness also changes who sells the model. Weighting by volume, an open-weight model is served by 14.9 independent hosting providers against 3.0 for a proprietary model, and only 4% of open-weight volume runs on a model with a single host, against 12% of proprietary volume (Figure 7 in the appendix). The price a user pays for an open model is set by hosts competing to serve the same weights rather than by the organisation that trained them, though among multi-host models the ninetyth-percentile host still charges about 1.6 times the cheapest.

4 Price setting and demand

4.1 Frequency of price revision

The archived catalogues let us follow the posted price of a model version from its first appearance. Between consecutive captures, a median of three days apart, the price changes 2.9% of the time, and 54% of models never change price at all. The left panel of Figure 3 shows 56% of models still at their launch price after six months and 45% after a year. Conditional on a change the median revision is a cut of about 5 log points, and 58% of changes are cuts.




Figure 3: Left, Kaplan-Meier survival of a model's launch price. Right, coefficients β_τ from (5) around 306 revisions, 95% intervals clustered by model.

Expressed the way the price-setting literature does, 12.6% of model-months contain a price change, an expected duration of about 7.4 months. That is at the low end of the range for consumer goods and comparable to online retail [Bils and Klenow, 2004, Nakamura and Steinsson, 2008, Gorodnichenko and Talavera, 2017, Cavallo, 2018], and that comparison is itself the surprise. The menu cost that explains rigidity in retail is close to zero when the price is a field in an API catalogue, and the managerial-attention frictions estimated for online sellers [Ellison et al., 2018] bind far less on a provider tracking a handful of models than on a retailer tracking thousands. What is missing is not the ability to reprice but the incentive. Over the few weeks a release event study spans, an incumbent's posted price is unchanged with probability near 0.9, so a design that looks for the competitive effect of open weights in incumbent price series is looking at a variable that rarely moves.

4.2 Volume response to price revisions

We use the rare revisions to estimate how volume responds. Merging the archived leaderboards with the archived catalogues gives 19,803 model-date observations on 634 models between November 2023 and May 2026, containing 708 observations that follow a price change. We regress log volume on the log fixed-weight price with model and date fixed effects, so that identification comes only from within-model revisions, net of anything common to the market on a date,

$$\log y_{jt} = \beta \log p_{jt} + \alpha_j + \gamma_t + \varepsilon_{jt}. \quad (4)$$

Here α_j is a model effect, γ_t a date effect and ε_{jt} the residual. We estimate $\hat{\beta} = -0.53$ (s.e. 0.10) for tokens and -0.46 (s.e. 0.09) for requests, with indistinguishable estimates for open-weight and proprietary models. Dropping the model fixed effects, so that cross-model variation identifies β , moves the estimate to -0.24 , which is what one would expect if higher-priced models carry unobserved quality. Figure 4 shows that the estimate survives every alternative we tried, holding between -0.66 and -0.46 across ten specifications and between -0.58 and -0.51 when models are dropped one at a time.




Figure 4: Own-price elasticity across specifications, with 95% intervals clustered by model. Shading marks the range spanned by dropping each of the 268 models whose price ever moves, one at a time.

The right panel of Figure 3 traces the dynamics around 306 revisions, restricted to the captures that report non-overlapping weekly buckets so that the path is not an artefact of a trailing aggregation

174 window. The specification replaces the level term in (4) by leads and lags interacted with the size of
 175 the revision Δ_{jk} ,

$$\log y_{jt} = \sum_{\tau \neq -1} \beta_{\tau} \Delta_{jk} \mathbf{1}\{t - T_{jk} = \tau\} + \delta_{jk} + \gamma_t + \varepsilon_{jt}, \quad (5)$$

176 where T_{jk} is the date of revision k to model j and δ_{jk} a revision-specific effect. Volume is flat before
 177 the change and settles about 0.3 to 0.45 log points lower per log point of price increase within a few
 178 weeks. The pre-period coefficients are small and individually insignificant, which is the relevant
 179 check. Had providers raised prices after volume rose, the estimate would be biased toward zero, and
 180 we find no such pattern.

181 An elasticity near one-half is not consistent with a static Bertrand equilibrium in posted prices, which
 182 would require demand to be elastic at the optimum. We read it as a short-run object. Switching a
 183 production workload to another model requires re-testing prompts, harnesses and evaluation suites,
 184 the kind of switching cost that sustains price dispersion in equilibrium [Klemperer, 1995, Farrell and
 185 Klemperer, 2007], so the relevant margin adjusts over months rather than days. That also rationalises
 186 the rigidity. A provider that would gain little volume from a price cut, and lose little from holding
 187 steady, has weak incentive to reprice a model already in the field.

188 5 Model entry

189 5.1 Decomposition of volume growth

190 The daily panel covers every model with usage for the four weeks to 20 August 2026, a window
 191 containing 10 releases by providers with priced, scored models, of which 5 are open-weight, and over
 192 which the market grew 65%. The left panel of Figure 5 decomposes daily volume into incumbents
 193 and models released during the window. Entrants reached 28% of platform tokens within four weeks,
 194 69% of it open-weight. Incumbent volume did not fall. It grew 18%. Models released during the
 195 window account for 72% of net growth over it, and incumbent expansion for the rest.




Figure 5: Left, daily platform volume split between models already serving traffic at the start of the window and models released during it. Right, event-time coefficients from (6) per unit of exposure (7), 95% intervals clustered by incumbent.

196 5.2 Incumbent responses to entry

197 Aggregate volume can hide reallocation, so we test for it directly. For each release k at date T_k we
 198 stack all incumbents j that were already serving traffic, and estimate

$$\log y_{jt} = \beta \mathbf{1}\{t \geq T_k\} \times \tau_{jk} + \delta_{jk} + \lambda_{kt} + X'_{jt} \theta + \varepsilon_{jkt}, \quad (6)$$

199 where δ_{jk} is an incumbent-by-entrant effect and λ_{kt} an entrant-by-day effect, so that β compares
 200 incumbents exposed to the same release on the same day. The exposure measure combines techno-
 201 logical proximity with the entrant's relative value for money,

$$\tau_{jk} = S_{jk} r_{jk}, \quad r_{jk} = \log \frac{Q_k}{p_k} - \log \frac{Q_j}{p_j}, \quad (7)$$

202 with S_{jk} from (2). Because capability and price both trend, X_{jt} also contains post-release interac-
 203 tions with the incumbent's own capability and price.

We find no diversion. The estimate of β is 0.01 (s.e. 0.07), and the event-time path in the right panel of Figure 5 is flat on both sides of the release. The confidence interval excludes differential declines larger than 0.13 log points per unit of exposure, tightening to 0.10 and 0.04 under the Euclidean and cosine measures in Table 4. Ten releases bound the effect rather than pin it down, but at this rate of growth the displacement channel is plainly small relative to expansion.

5.3 Longer-horizon estimates

The archived panel lets us look further out. We take the 19 open-weight releases that reached at least one percent of platform tokens and compare aggregate volume of models already in the market over the eight weekly observations before and after each. Incumbent volume fell in 32% of these episodes, and the median change was +0.25 log points against a median market change of +0.52. DeepSeek R1 [DeepSeek-AI, 2025] is the sharpest case, with platform volume rising 0.56 log points over the sixteen weeks around it while incumbent volume held flat at -0.03 . Incumbents lose share because the market grows around them, not because they lose volume.

6 Demand system and counterfactuals

6.1 Specification

Because posted prices do not respond and incumbents are not displaced, the value of open-weight competition is measured on the choice set rather than on prices. We estimate a nested-logit demand system over the 172 models with prices and capability scores, nesting by capability tercile so that models of similar capability are closer substitutes. Writing $g(j)$ for the nest of model j and $\sigma \in [0, 1)$ for the within-nest correlation [Cardell, 1997, Berry, 1994], shares take the form

$$s_j = \frac{\exp(\delta_j/(1-\sigma))}{D_{g(j)}} \cdot \frac{D_{g(j)}^{1-\sigma}}{1 + \sum_h D_h^{1-\sigma}}, \quad D_g = \sum_{k \in g} \exp(\delta_k/(1-\sigma)), \quad (8)$$

with mean utility

$$\delta_j = \beta_p \log p_j + \alpha' x_j + \xi_j, \quad (9)$$

where x_j collects the capability index, context length, modality, reasoning support, latency and age, α the coefficients on them and ξ_j the unobserved quality of model j . The outside option is the unscored fringe, whose share s_0 is 2.4% of tokens, and $s_{j|g(j)}$ denotes model j 's share within its own nest. Shares invert to mean utilities [Berry, 1994],

$$\delta_j = \log s_j - \log s_0 - \sigma \log s_{j|g(j)}, \quad (10)$$

so the system reproduces the data exactly and the counterfactual does not turn on how well the characteristics fit. Consumer surplus follows from the expected maximum utility [Small and Rosen, 1981, McFadden, 1974],

$$CS = \frac{1}{|\beta_p|} \log \left(1 + \sum_h D_h^{1-\sigma} \right). \quad (11)$$

Deleting the open-weight models changes the denominators in (8), redistributing their volume across the remaining models and the outside option, so the counterfactual depends only on (β_p, σ) .

6.2 Parameter choice

Cross-sectional share variation does not identify β_p and σ here. The differentiation instruments of Gandhi and Houde [2019] are functions of the distribution of rival characteristics, which in this market is a distribution over an essentially one-dimensional index, so they are near-collinear with a model's own capability and with its unobserved quality. Parameter count, the obvious cost shifter, correlates 0.65 with the capability index and moves the price coefficient from -0.79 under least squares to -0.27 despite a first-stage F above 900, a violated exclusion restriction rather than a corrected bias.

We therefore fix $\beta_p = -1$ and report every result across a range of σ [Nevo, 2000, Conlon and Gortmaker, 2020]. Both choices are conservative. The price coefficient exceeds in magnitude the within-model estimate of Section 4.2 and the long-run elasticity of about -0.6 implied by (12),

and because surplus enters (11) divided by $|\beta_p|$ a larger magnitude yields a smaller dollar figure. Finding no differential loss within the entrant’s own capability tier is what σ near zero predicts, and the surplus loss is largest there, so $\sigma = 0.4$ is conservative too.

6.3 Removing open-weight models

Removing open-weight models from the choice set at posted prices, the average price paid per token rises from \$0.70 to \$1.80, an increase of 156%. Consumer surplus (11) falls 21%, worth about \$2.4m per day at current volume. The percentage change does not depend on β_p , which enters the baseline and the counterfactual identically, and only the conversion into dollars scales with it. The loss runs from \$3.9m to \$1.3m per day as σ moves from 0 to 0.7 (Table 5), and raising the outside share from the observed 2.4% to 50% raises the proportional loss from 21% to 47% while lowering the dollar figure to \$1.0m per day. The price effect is invariant to both, because it is governed by the composition of the surviving choice set rather than by the strength of substitution. Concentration moves against the presumption that open weights fragment the market. the provider Herfindahl index on tokens rises from 897 to 1048, because open-weight volume is spread across more suppliers than the proprietary volume that replaces it.




Figure 6: Left, for each of the 29 task markets, the price increase implied by removing open-weight models against their share of task spending, marker area proportional to the task’s share of platform spending. Right, aggregate surplus loss per day across σ and β_p .

These numbers hold prices fixed, which given the rigidity of Section 4.1 makes the exercise a statement about variety in the tradition of Trajtenberg [1989], Hausman [1996] and Petrin [2002]. Allowing providers to reprice would add to the loss, so the estimates are a lower bound.

The left panel of Figure 6 repeats the exercise separately in each of the 29 task-level markets, where shares of spending are reported directly. The implied price increase rises almost one for one with the open-weight share of task spending, with a slope of 0.98. Weighted by spending, prices rise 18% and surplus falls 13%, with the largest effects in agentic and coding work, where open-weight models hold about a third of spending, and the smallest in data extraction and transformation, where they hold under a sixth. These are smaller than the aggregate because task shares are measured on spending, which weights the expensive proprietary models that open weights do not displace. Table 2 groups the tasks into families. The effect tracks open-weight penetration closely, reaching 20% in coding work and 7% in data extraction and transformation.

Table 2: The open-weight counterfactual inside each family of task markets, weighted by each task’s share of platform spending. Shares and changes are averages over the tasks in the family.

Task family	Tasks	Share of spending	Open-weight share	Price without open
Agent	5	0.29	0.33	+19%
Code	9	0.30	0.31	+20%
General	12	0.22	0.25	+19%
Data	2	0.09	0.15	+7%

7 Interpretation

7.1 An accounting framework

The three results fit together through the share of demand that is in play. Let M_t be total demand growing at rate g per period, and let a fraction ϕ of existing demand be revisited each period as workloads are rebuilt or re-evaluated. Contestable demand is $(g + \phi)M_t$, and a model’s volume evolves as

$$y_{j,t+1} = (1 - \phi) y_{jt} + \nu_j (g + \phi) M_t, \quad (12)$$

where ν_j is its share of contestable demand. An incumbent loses volume only when $\nu_j(g + \phi)M_t < \phi y_{jt}$, so with $g \approx 0.13$ per week in our window an incumbent holding an average share of new demand keeps growing. An entrant reaches a share of about $\nu_k(g + \phi)/(1 + g)$ after one period, taking volume quickly without anyone shrinking. And a permanent price change reaches only the contestable margin, so the response measured τ periods out is attenuated by roughly $1 - (1 - \phi - g)^\tau$ relative to its long-run value.

That last point bounds how much of the inelasticity is an artefact of slow adjustment. With $\phi + g \approx 0.23$ per week and the six weekly observations over which Figure 3 shows the response settling, the attenuation factor is about 0.8 and the implied long-run elasticity is near -0.6 . Demand is inelastic even after adjustment, so the first-order condition of a provider maximising current profit is not satisfied at the posted price. Either providers price against a margin we do not observe, or the posted price is not chosen to maximise current profit at all, and the rigidity of Section 4.1 favours the second.

The foregone savings Nagle and Yue [2025] identify are a puzzle if demand adjusts quickly and are exactly what one should expect when the short-run elasticity is about one-half and switching a workload is costly. They accrue as workloads turn over rather than as incumbents reprice. Our price-of-capability gaps are smaller than theirs, consistent with proprietary providers having released cheaper tiers since and with the rebound in the open-weight volume share in Figure 2.

7.2 Limitations

Three qualifications bound what the estimates cover. The market is one router that selects for price-sensitive developers, so the estimates describe the segment where substitution toward open weights is easiest. The demand parameters are calibrated rather than estimated, which is why every counterfactual is reported over a range and at conservative values. And four weeks of releases support a bound on diversion rather than a point estimate, since the staggered-entry estimators that handle heterogeneous treatment timing [Callaway and Sant’Anna, 2021, de Chaisemartin and D’Haultfœuille, 2020, Sun and Abraham, 2021, Borusyak et al., 2024] need more events than the window provides.

8 Conclusion

Open-weight models constrain proprietary providers on a margin the usual measurement misses. Posted prices behave like a fixed characteristic of a model version. More than half of models are never repriced, and the monthly frequency of a change is comparable to groceries in a setting where changing the number costs nothing. Demand explains why. The within-model elasticity is about one-half, and correcting for slow reallocation moves it only to -0.6 , so a provider gains little volume from cutting the price of a model already in the field and does not cut it. Entry is absorbed by growth. Over four weeks, models released inside the window took more than a quarter of platform tokens while incumbent volume rose, and we find no differential loss for the incumbents nearest an entrant. Across nineteen major open-weight releases since 2024, incumbent volume fell in fewer than a third of episodes.

The constraint is nonetheless large. Removing open-weight models from an estimated demand system raises the average price paid per token by more than half again and costs users roughly a fifth of their surplus, on the order of a million dollars a day on one platform, with the largest effects in agentic and coding work. It also raises provider concentration rather than lowering it, because open-weight volume is spread across more suppliers than the proprietary volume that would replace it. An assessment built on incumbent price responses will therefore find nothing, not because the constraint is absent but because it operates on the price at which the next unit of demand is served.

References

- Y. Bengio et al. International AI Safety Report 2026. arXiv:2602.21012, 2026.
- S. Berry. Estimating discrete-choice models of product differentiation. *RAND Journal of Economics*, 25(2):242–262, 1994.
- S. Berry and P. Haile. Identification in differentiated products markets using market level data. *Econometrica*, 82(5):1749–1797, 2014.
- S. Berry, J. Levinsohn, and A. Pakes. Automobile prices in market equilibrium. *Econometrica*, 63(4):841–890, 1995.
- M. Bils and P. J. Klenow. Some evidence on the importance of sticky prices. *Journal of Political Economy*, 112(5):947–985, 2004.
- R. Bommasani et al. On the opportunities and risks of foundation models. arXiv:2108.07258, 2021.
- K. Borusyak, X. Jaravel, and J. Spiess. Revisiting event-study designs: Robust and efficient estimation. *Review of Economic Studies*, 91(6):3253–3285, 2024.
- T. F. Bresnahan. Empirical studies of industries with market power. In *Handbook of Industrial Organization*, volume 2, pages 1011–1057. North-Holland, 1989.
- B. Callaway and P. H. C. Sant’Anna. Difference-in-differences with multiple time periods. *Journal of Econometrics*, 225(2):200–230, 2021.
- N. S. Cardell. Variance components structures for the extreme-value and logistic distributions with application to models of heterogeneity. *Econometric Theory*, 13(2):185–213, 1997.
- A. Cavallo. Scraped data and sticky prices. *Review of Economics and Statistics*, 100(1):105–119, 2018.
- A. Cavallo and R. Rigobon. The billion prices project: Using online prices for measurement and research. *Journal of Economic Perspectives*, 30(2):151–178, 2016.
- W.-L. Chiang et al. Chatbot Arena: An open platform for evaluating LLMs by human preference. In *International Conference on Machine Learning*, 2024.
- C. Conlon and J. Gortmaker. Best practices for differentiated products demand estimation with PyBLP. *RAND Journal of Economics*, 51(4):1108–1161, 2020.
- C. de Chaisemartin and X. D’Haultfoeulle. Two-way fixed effects estimators with heterogeneous treatment effects. *American Economic Review*, 110(9):2964–2996, 2020.
- DeepSeek-AI. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. arXiv:2501.12948, 2025.
- M. Demirer, A. Fradkin, N. Tadelis, and S. Peng. The emerging market for intelligence: Pricing, supply, and demand for LLMs. NBER Working Paper 34608, 2025.
- G. Ellison, C. Snyder, and H. Zhang. Costs of managerial attention and activity as a source of sticky prices: Structural estimates from an online market. NBER Working Paper 24680, 2018.
- J. Farrell and P. Klemperer. Coordination and lock-in: Competition with switching costs and network effects. In *Handbook of Industrial Organization*, volume 3, pages 1967–2072. North-Holland, 2007.
- A. Gandhi and J.-F. Houde. Measuring substitution patterns in differentiated-products industries. NBER Working Paper 26375, 2019.
- Y. Gorodnichenko and O. Talavera. Price setting in online markets: Basic facts, international comparisons, and cross-border integration. *American Economic Review*, 107(1):249–282, 2017.
- S. Greenstein and F. Nagle. Digital dark matter and the economic contribution of Apache. *Research Policy*, 43(4):623–631, 2014.

- 366 J. A. Hausman. Valuation of new goods under perfect and imperfect competition. In *The Economics*
367 *of New Goods*, pages 207–248. University of Chicago Press, 1996.
- 368 J. Hoffmann et al. Training compute-optimal large language models. In *Advances in Neural Infor-*
369 *mation Processing Systems*, 2022.
- 370 J. Kaplan et al. Scaling laws for neural language models. arXiv:2001.08361, 2020.
- 371 S. Kapoor, R. Bommasani, K. Klyman, et al. Position: On the societal impact of open foundation
372 models. In *International Conference on Machine Learning*, 2024.
- 373 P. Klemperer. Competition when consumers have switching costs. *Review of Economic Studies*,
374 62(4):515–539, 1995.
- 375 P. J. Klenow and B. A. Malin. Microeconomic evidence on price-setting. In *Handbook of Monetary*
376 *Economics*, volume 3, pages 231–284. North-Holland, 2011.
- 377 A. Korinek and J. Vipra. Concentrating intelligence: Scaling and market structure in artificial intel-
378 ligence. *Economic Policy*, 40(121):225–256, 2025.
- 379 J. Lerner and J. Tirole. Some simple economics of open source. *Journal of Industrial Economics*,
380 50(2):197–234, 2002.
- 381 P. Liang et al. Holistic evaluation of language models. *Transactions on Machine Learning Research*,
382 2023.
- 383 D. McFadden. Conditional logit analysis of qualitative choice behavior. In *Frontiers in Economet-*
384 *rics*, pages 105–142. Academic Press, 1974.
- 385 F. Nagle. Open source software and firm productivity. *Management Science*, 65(3):1191–1215,
386 2019.
- 387 F. Nagle and D. Yue. The latent role of open models in the AI economy. Working paper, 2025.
- 388 E. Nakamura and J. Steinsson. Five facts about prices: A reevaluation of menu cost models. *Quar-*
389 *terly Journal of Economics*, 123(4):1415–1464, 2008.
- 390 A. Nevo. A practitioner’s guide to estimation of random-coefficients logit models of demand. *Jour-*
391 *nal of Economics and Management Strategy*, 9(4):513–548, 2000.
- 392 A. Nevo. Measuring market power in the ready-to-eat cereal industry. *Econometrica*, 69(2):307–342,
393 2001.
- 394 National Telecommunications and Information Administration. Dual-use foundation models with
395 widely available model weights. U.S. Department of Commerce report, 2024.
- 396 A. Petrin. Quantifying the benefits of new products: The case of the minivan. *Journal of Political*
397 *Economy*, 110(4):705–729, 2002.
- 398 M. Rysman. The economics of two-sided markets. *Journal of Economic Perspectives*, 23(3):125–
399 143, 2009.
- 400 K. A. Small and H. S. Rosen. Applied welfare economics with discrete choice models. *Economet-*
401 *rica*, 49(1):105–130, 1981.
- 402 L. Sun and S. Abraham. Estimating dynamic treatment effects in event studies with heterogeneous
403 treatment effects. *Journal of Econometrics*, 225(2):175–199, 2021.
- 404 M. Trajtenberg. The welfare analysis of product innovations, with an application to computed to-
405 mography scanners. *Journal of Political Economy*, 97(2):444–479, 1989.

A Data construction

Endpoints. The catalogue endpoint returns, for every listed model, the posted input and output price per token, context length, input and output modalities, a creation timestamp, and the Hugging Face repository identifier when the weights are public. The leaderboard endpoint returns tokens, requests, tool calls and media counts by model and serving variant. A benchmark endpoint returns third-party capability scores and usage-weighted effective prices. A task endpoint returns, for each of 29 usage categories, the category’s share of platform spending and the share held by each of its ten largest models, in both spending and token terms. An endpoint listing gives each hosting provider’s price, quantisation and uptime for a given model. All are public and unauthenticated.

Daily usage. Each model’s public page embeds its own trailing-month daily usage series in the page’s streamed render payload. We request one page per model with observed usage and parse the payload, which yields 457 models over 30 days. The same payload carries per-endpoint latency and throughput percentiles, which we aggregate to the model level by request count.

Historical coverage. Internet Archive captures supply history. Captures of the catalogue endpoint give a price and characteristic panel from July 2023. Captures of the leaderboard page embed the leaderboard payload in the same streamed format, so we recover model-level usage without relying on the rendered page. Two vintages are present and we treat them separately: captures before late 2024 embed a series of non-overlapping weekly buckets, and later captures embed a single trailing-window snapshot. Shares are comparable across vintages; levels are not, so all specifications include date fixed effects, and Table 3 reports the two vintages separately. We fetch captures at a fixed spacing with adaptive backoff, which is the binding cost of the collection: the archive requests take a few hours, and every other step runs in minutes on a laptop.

Merging. Usage and price observations are matched by model identifier and nearest date within ten days. Openness is taken from the current catalogue where the model is still listed, because the weight-repository field is absent from the earliest captures. Prices are dollars per million tokens throughout.

B Capability imputation

Benchmark scores are missing for smaller models. For each missing dimension we fit a ridge regression on the dimensions the model does have, using models scored on both, and predict the missing value. Masking half of the fully scored models and predicting the held-out values gives R^2 of 0.99, 0.97 and 0.96 for the reasoning, coding and agentic dimensions, with mean absolute errors of 1.7, 5.0 and 3.1 points against standard deviations of 16.4, 22.6 and 18.5. Imputation raises coverage from 109 models and 77% of tokens to 180 models and 91% of tokens. All results in the main text hold on the directly scored subsample.

C Robustness

Table 3 reports the elasticity across price measures, outcome measures, capture vintages, openness and sample period. The estimate stays between -0.38 and -0.66 . The realised blended price, which mixes posted prices with the model’s own input-output composition, gives the largest magnitude, as expected if a shift toward prompt-heavy workloads lowers the average price per token and raises volume at the same time; the fixed-weight index removes that channel and is our preferred measure.

The estimate is not driven by any one model: dropping each of the 268 models whose price ever moves, one at a time, leaves it between -0.58 and -0.51 . Restricting to models holding at least 0.01% of tokens gives -0.378 (s.e. 0.127).

Two further checks bear on the elasticity. A price cut may accompany the launch of a replacement by the same provider, which would depress the old model’s volume for reasons other than price. Adding an indicator for observations within 28 days of a same-provider release, which covers 38% of the panel, leaves the estimate at -0.534 ; dropping those observations gives -0.499 (s.e. 0.123).

Submitted to 40th Conference on Neural Information Processing Systems (NeurIPS 2026). Do not distribute.

Table 4 reports the diversion test under three definitions of technological distance. None of the six estimates is distinguishable from zero. The Euclidean and cosine measures give tighter standard errors than the Mahalanobis measure, and their confidence intervals bound the differential decline for the most exposed incumbents at 0.10 and 0.04 log points respectively.

Table 3: Own-price elasticity across specifications. Each row is a separate estimate of (4) with model and date fixed effects; standard errors clustered by model.

Specification	Estimate	Std. error	Observations	Models
Fixed-weight price (baseline)	−0.534	0.102	19,803	634
Input price	−0.526	0.096	19,803	634
Output price	−0.485	0.109	19,803	634
Realised blended price	−0.658	0.095	19,803	634
Requests rather than tokens	−0.463	0.087	19,803	634
Weekly captures only	−0.472	0.123	15,133	432
Trailing captures only	−0.476	0.125	4,670	428
Open-weight models	−0.582	0.116	8,873	321
Proprietary models	−0.534	0.204	10,930	316
2025 onwards	−0.527	0.100	7,906	542

Table 4: Diversion at entry under alternative distance measures. Each row estimates (6) with incumbent-by-entrant and entrant-by-day fixed effects and post-release controls for the incumbent’s own capability and price.

Distance	Exposure	Estimate	Std. error	Observations
Mahalanobis	similarity $\times$ relative value	0.010	0.070	11,820
Mahalanobis	similarity	0.019	0.120	11,820
Euclidean	similarity $\times$ relative value	−0.006	0.047	11,820
Euclidean	similarity	0.008	0.065	11,820
Cosine	similarity $\times$ relative value	0.022	0.031	11,820
Cosine	similarity	0.083	0.086	11,820




Figure 7: Left: distribution of tokens across the number of independent hosting providers serving the model. Right: cumulative distribution, across models with more than one host, of the ratio between the ninetieth-percentile and the cheapest posted host price.

D How good and how cheap an open model has to be

The counterfactual removes the open-weight models that exist. The opposite exercise asks what a new one would have to look like to matter. We insert a single open-weight entrant with capability $(1 - \rho) \max_k Q_k$ and price p , assign it the mean utility implied by the fitted gradient of (9) in capability and log price, and recompute shares. Figure 8 traces the effect on proprietary volume. The iso-effect lines are steep, which is the substantive point: price moves the outcome far more than the last few points of capability. An entrant that matches the frontier exactly but is priced at \$0.50 per million tokens takes 6.6% of platform tokens and cuts proprietary volume by 7%. An entrant a tenth below the frontier but priced at \$0.05 takes 17% and cuts proprietary volume by 18%. At the frontier’s own price, an open model that matches it takes essentially nothing.




Figure 8: Change in proprietary token share when one hypothetical open-weight model is added to the choice set, by its price and its capability shortfall against the frontier, at $\beta_p = -1$ and $\sigma = 0.4$. Contours mark 2, 5 and 10 percent losses.

E Counterfactual sensitivity

Table 5 reports the counterfactual across the within-tier substitution parameter, and robustness. json reports it across the size of the outside option as well.

Table 5: The open-weight counterfactual across substitution parameters, at $\beta_p = -1$. Prices are dollars per million tokens; the surplus loss is evaluated at the platform’s volume on the reference day.

σ	Price paid (\$/M)	Price without open (\$/M)	Surplus change (%)	Surplus loss (\$m per day)	HHI (without open)
0.0	0.70	1.82	-33.8	3.9	1109
0.2	0.70	1.81	-27.5	3.2	1082
0.4	0.70	1.80	-21.0	2.4	1048
0.6	0.70	1.79	-14.3	1.7	1006
0.7	0.70	1.78	-10.8	1.3	982

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and Section 1 state the three measurement results and the counterfactual, each of which is established in Sections 3–5, and both state that the market studied is one inference router.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Section 6 gives a limitations paragraph covering market coverage, the capability index, the calibrated demand parameters and the power of the diversion test.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper contains no theorems. The assumptions of the demand model are stated in Section 5.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Every input is a public endpoint or an Internet Archive capture; Appendix A documents the collection and construction, and the released code reproduces the panels and all reported numbers.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The collection and analysis code is released, and the raw inputs can be rebuilt from the public endpoints and archive captures the code queries.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Sample construction, variable definitions, fixed effects and clustering are given in Sections 3–5 and Appendix A.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: All regressions report standard errors clustered by model, and the calibrated demand parameters are reported across a grid rather than at a single value.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.

- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: All estimation runs in a few minutes on a laptop; the only substantial cost is the rate-limited archive collection, which is documented in Appendix A.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research conforms with the NeurIPS Code of Ethics. It uses aggregate, model-level public data and no personal data.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Section 6 discusses the implication for how the competitive effect of open weights should be assessed.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper releases analysis code and no models, weights or scraped personal data.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The data sources are public API and web endpoints and Internet Archive captures; Appendix A records the sources and access terms.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.

- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The released code is documented, and Appendix A describes the derived panels it produces.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The paper does not involve crowdsourcing or human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The paper does not involve human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.

- 784 • Depending on the country in which research is conducted, IRB approval (or equivalent)
785 may be required for any human subjects research. If you obtained IRB approval,
786 you should clearly state this in the paper.
- 787 • We recognize that the procedures for this may vary significantly between institutions
788 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
789 guidelines for their institution.
- 790 • For initial submissions, do not include any information that would break anonymity
791 (if applicable), such as the institution conducting the review.

792 16. Declaration of LLM usage

793 Question: Does the paper describe the usage of LLMs if it is an important, original, or
794 non-standard component of the core methods in this research? Note that if the LLM is used
795 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
796 scientific rigor, or originality of the research, declaration is not required.

797 Answer: [Yes]

798 Justification: Large language models were used for coding assistance and for editing prose,
799 not as part of the research method. The data collection, estimation and reported results
800 were specified and checked by the authors.

801 Guidelines:

- 802 • The answer [N/A] means that the core method development in this research does not
803 involve LLMs as any important, original, or non-standard components.
- 804 • Please refer to our LLM policy in the NeurIPS handbook for what should or should
805 not be described.