
Entry Without Displacement: How Open-Weight Models Compete in the Market for Inference

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Open-weight models are widely expected to constrain the pricing power of propri-
2 etary model providers, but the mechanism has not been measured. We assemble a
3 market-wide panel of the models served by a large inference router: posted prices,
4 capability benchmarks, weight-release status and revealed usage for 457 models,
5 extended back to 2023 with archived catalogues. Three facts emerge. Competition
6 works through releases rather than repricing: 54% of models are never repriced
7 at all, and a posted price changes in only 12.6% of model-months, even though
8 changing it costs a provider nothing. When prices do change, volume responds
9 weakly, with an own-price elasticity of -0.53 (s.e. 0.10) identified from 708 re-
10 visions with model and date fixed effects. Entry does not displace incumbents:
11 over four weeks, models released inside the window took 28% of tokens while
12 incumbent volume grew 18%, and the confidence interval rules out differential
13 losses above 0.13 log points for the incumbents closest to an entrant in capability
14 space. Removing open-weight models from the choice set in an estimated nested-
15 logit demand system raises the average price paid per token by 156% and costs
16 users about \$2.4m per day on this platform, while raising provider concentration.
17 Open weights discipline the price at which new demand is served, not the prices
18 incumbents post.

19 1 Introduction

20 Whether open-weight models constrain the market power of proprietary model providers is now a
21 policy question as much as an economic one. Reports commissioned by governments name it as an
22 open evidence gap [Bengio et al., 2026], and the theoretical case cuts both ways: scale economies in
23 training push toward concentration [Korinek and Vipra, 2025], while freely downloadable weights
24 supply a competitive fringe that any hosting provider can serve. Recent empirical work has begun
25 to describe the market. Demiret et al. [2025] document a collapse in the price of a unit of measured
26 intelligence and a proliferation of suppliers; Nagle and Yue [2025] show that a large share of propri-
27 etary usage is directed at models that are dominated on price and measured capability by an open
28 alternative, and estimate the savings foregone.

29 Both of those results are about the allocation of demand across models at a point in time. Neither
30 identifies the channel through which open weights would constrain proprietary pricing, and the
31 two natural channels have opposite implications for policy. If open weights force incumbents to
32 cut posted prices, then the constraint is visible in the price series of established models and its
33 strength can be read off their price responses. If instead open weights capture new demand while
34 incumbents hold their prices, the constraint operates on the terms at which the market grows, and
35 any measurement based on incumbent price responses will find nothing.

We assemble a panel that can separate the two. Our data come from a large inference router that publishes model-level usage, prices and capability scores. We combine the current cross-section (457 models with observed usage, 175 of them open-weight), a daily usage panel covering all of them for one month, task-level market shares for 29 distinct uses, per-provider hosting terms, and a panel of archived price catalogues and usage leaderboards reaching back to 2023. The panel used for price responses contains 19,803 model-date observations on 634 models.

Three facts follow, and they point the same way. First, the posted price of a given model version behaves like a fixed characteristic of the product: 54% of models are never repriced, and the monthly frequency of a price change is 12.6%, comparable to consumer goods [Nakamura and Steinsson, 2008, Cavallo, 2018] in a setting where the menu cost that explains rigidity there does not exist. Second, when prices do change, volume moves little. Exploiting 708 revisions with model and date fixed effects, we estimate an own-price elasticity of -0.53 (s.e. 0.10), below the value near one reported by Demirer et al. [2025] from cross-model variation and far below what homogeneous-goods reasoning would imply. Third, entry is absorbed by growth rather than by displacement. Over the four weeks of the daily panel the market grew 65%; models released inside the window took 28% of tokens, of which 5 of 10 releases were open-weight, and yet incumbent volume rose 18%. In an incumbent-by-entrant event study we find no differential decline for the incumbents technologically closest to an entrant, with a confidence interval that excludes declines larger than 0.13 log points.

Read together, these facts say that the competitive constraint from open weights does not show up where one would look for it. We therefore quantify it where it does show up. We estimate a nested-logit demand system over the 172 models with prices and capability scores, invert it to recover mean utilities that reproduce observed shares, and delete open-weight models from the choice set. Users would pay 156% more per token, and the surplus loss is worth about \$2.4m per day on this platform alone. The loss is not evenly spread: across 29 task-level markets, the implied price increase rises almost one for one with the open-weight share of task spending, from 18% on average to more than twice that in agentic and coding work. Concentration moves the other way from the usual presumption: removing open weights raises the provider Herfindahl index from 897 to 1048.

Our contribution is to identify the margin on which open-weight competition operates and to price it. The measurement of price rigidity and of the within-model demand response is, to our knowledge, new to this market, and it explains why the reallocation gains documented by Nagle and Yue [2025] persist: with inelastic short-run demand, a provider has little to gain from cutting the price of a model already in the field, and users have little reason to move quickly. We also release the collection code, which reconstructs the entire panel from public endpoints.

2 Data

Sources. We use the public API and web endpoints of OpenRouter, a router that serves requests to hundreds of models from many providers through one interface. Four endpoints supply the raw material: a catalogue giving posted input and output prices, context length, release timestamp and the Hugging Face repository for models whose weights are published; a leaderboard giving tokens, requests and tool calls by model; a benchmark table reporting third-party capability scores; and a task table giving each model’s share of spending in 29 usage categories over a 30-day window. Each model’s own page embeds a daily usage series for the trailing month, which we collect for every model with observed usage. We take per-provider hosting terms from the endpoint listing, and parameter counts from Hugging Face. Historical coverage comes from Internet Archive captures of the catalogue and of the leaderboard; the latter embed the leaderboard payload in the page’s streamed data, which we parse directly. Appendix A documents each step.

Openness. We classify a model as open-weight if the catalogue records a public weight repository for it. The classification is model-specific rather than provider-specific, which matters here: several providers, including Google, OpenAI and Alibaba, ship both open-weight and proprietary models, and a provider-level rule would misclassify them.

Capability. Capability enters as a vector of third-party scores covering reasoning, coding and agentic use, supplemented by pairwise-comparison scores from a design-oriented arena. These measures are highly collinear. The first principal component of the three benchmark dimensions explains 98% of their variance, and 91% when the two arena blocks are added, so we summarise capability

Table 1: The estimation sample, measured on the last complete day of the daily panel. Prices are dollars per million tokens for the blended input-output mix. Capability is the first principal component of the benchmark block. Shares are within the estimation sample, which covers 91% of platform tokens.

	Models	Median price	Median capability	Median context (thousands)	Token share	Spend share
Open-weight	86	0.29	30.5	256	0.73	0.32
Proprietary	86	1.50	36.6	400	0.27	0.68
All scored	172	0.59	33.4	262	1.00	1.00

by that first component and report the underlying dimensions in the appendix. Scores are missing for smaller models; we impute a missing dimension by ridge regression on the dimensions a model does have, which recovers held-out values with R^2 above 0.96. The resulting sample of 180 scored models accounts for 91% of tokens.

Prices. Providers post separate input and output prices. Because the input-output mix differs sharply across models, we work with a blended price that weights the two by the observed token mix, and with a fixed-weight index that holds the mix at its sample mean when we need a price series free of feedback from usage. Table 1 summarises the estimation sample.

What the market is. The platform is one segment of the market for inference. It excludes traffic sent directly to a provider’s own API and consumer subscription products, and its users are developers who have chosen a router, which selects for price sensitivity and for willingness to switch. Our estimates therefore describe the segment in which open-weight substitution is easiest. If anything this biases us toward finding a strong competitive constraint, which sharpens the interpretation of the weak responses we report below.

3 Three facts

3.1 Capability is cheap, and open weights are cheaper

Figure 1 places every model with usage on the price-capability plane. Two features stand out. At any capability level below the frontier, prices differ by one to two orders of magnitude, so posted price is far from a sufficient statistic for what users buy. Second, the open-weight side of the market is systematically cheaper. Matching each proprietary model to the open-weight models within two points of it on the capability index, the median proprietary model costs 3.3 times the median open alternative, or 1.7 times when the comparison is weighted by volume. Averaged over actual usage, a token from an open-weight model costs \$0.35 per million against \$1.95 for a proprietary one. These gaps are smaller than the ninety percent difference reported by Demirer et al. [2025] and Nagle and Yue [2025] for 2025 data, which is what one would expect if proprietary providers have since released cheaper tiers; the frontier comparison drawn in Figure 1, which uses only the cheapest option available at each capability level, narrows the gap further to 1.4 times.

The right panel tracks the best capability obtainable from each side over time. The open-weight frontier follows the proprietary frontier with a median lag of 105 days. The gap is a lag, not a level: at the very top of the capability distribution proprietary models are unmatched at every date, but the set of tasks for which that matters shrinks as the open frontier advances.

The open-weight share of volume is also far from stable. In the archived leaderboards, which are internally consistent over time though not directly comparable in level with the daily panel, open-weight models held 86% of tokens in late 2023, fell to 29% by mid-2025, and recovered to 48% by mid-2026. The trough coincides with the sample used by Nagle and Yue [2025], whose finding that proprietary models dominate usage describes a market that has since moved.

Openness also changes the supply side. Weighting by volume, an open-weight model is served by 14.9 independent hosting providers against 3.0 for a proprietary model, and the price a user pays for an open model is set by hosts competing to serve the same weights rather than by the organisation that trained it. Open-weight models account for 68% of tokens but only 30% of spending, and the




Figure 1: The price of capability. Left: posted blended price against the capability index, with marker area proportional to the square root of token volume; the step functions give the cheapest price available at or above each capability level for each side of the market. Right: the highest capability available from each side by release date, with the shaded region marking the open-weight lag.

129 difference shows up in concentration: the provider Herfindahl index is 1183 measured on tokens and
 130 2439 measured on spending. By volume the market looks competitive; by revenue it does not.

131 3.2 Posted prices rarely move

132 The archived catalogues let us follow the posted price of a model version from its first appearance.
 133 Between consecutive captures, a median of three days apart, the price changes 2.9% of the time.
 134 Over a model’s observed life, 54% of models never change price at all. The left panel of Figure 2
 135 shows the survival of the launch price: 56% of models are still at their launch price after six months
 136 and 45% after a year. Conditional on a change, the median revision is a cut of about 5 log points,
 137 and 58% of changes are cuts.

138 Expressed the way the price-setting literature does, 12.6% of model-months contain a price change,
 139 an expected duration of about 7.4 months. That is at the low end of the range reported for consumer
 140 goods and comparable to online retail [Bils and Klenow, 2004, Nakamura and Steinsson, 2008,
 141 Cavallo, 2018], which is itself the surprise: the menu cost that explains rigidity in those settings, the
 142 physical and managerial cost of changing a posted price, is close to zero when the price is a field in an
 143 API catalogue. The distribution is also more polarised than in retail, where most goods are repriced
 144 eventually. Here more than half of models are never repriced at all. What is missing is not the
 145 ability to reprice but the incentive, which the next fact supplies. The consequence for measurement
 146 is immediate: over the few weeks that a release event study typically spans, an incumbent’s posted
 147 price is unchanged with probability near 0.9, so a design that looks for the competitive effect of open
 148 weights in incumbent price series is looking at a variable that rarely moves.

149 3.3 Demand for a given model is price-inelastic in the short run

150 We use the rare revisions to estimate how volume responds. Merging the archived leaderboards with
 151 the archived catalogues gives 19,803 model-date observations on 634 models between November
 152 2023 and May 2026, containing 708 observations that follow a price change. We regress log volume
 153 on the log fixed-weight price with model and date fixed effects, so identification comes only from
 154 within-model revisions, net of anything common to the market on a date:

$$\log y_{jt} = \beta \log p_{jt} + \alpha_j + \gamma_t + \varepsilon_{jt}. \quad (1)$$

155 Table 2 reports $\hat{\beta} = -0.53$ (s.e. 0.10) for tokens and -0.46 (s.e. 0.09) for requests. Open-weight
 156 and proprietary models give indistinguishable estimates. Dropping the model fixed effects, so that
 157 cross-model variation identifies β , moves the estimate to -0.24 , which is what one would expect if
 158 higher-priced models carry unobserved quality.

159 The right panel of Figure 2 traces the dynamics around 306 revisions, restricted to the captures that
 160 report non-overlapping weekly buckets so that the path is not an artefact of a trailing aggregation




Figure 2: Posted prices rarely move, and volume responds weakly when they do. Left: Kaplan-Meier survival of a model’s launch price in the archived catalogues. Right: coefficients from a distributed-lag version of (1) around 306 price revisions, scaled by the size of the revision, with model-by-event and date fixed effects and 95% intervals clustered by model.

Table 2: Own-price elasticity of demand. Each column is a separate regression of (1). Standard errors clustered by model in parentheses.

	Model and date FE		By openness		Date FE only
	Tokens	Requests	Open	Proprietary	Tokens
log price	−0.534 (0.102)	−0.463 (0.087)	−0.582 (0.116)	−0.534 (0.204)	−0.239 (0.077)
Observations	19,803	19,803	8,873	10,930	19,803
Models	634	634	321	316	634

161 window. Volume is flat before the change and settles about 0.3 to 0.45 log points lower per log
162 point of price increase within a few weeks. The pre-period coefficients are small and individually
163 insignificant, which is the relevant check: if providers raised prices after volume rose, the estimate
164 would be biased toward zero, and we find no such pattern.

165 An elasticity near one-half is not consistent with a static Bertrand equilibrium in posted prices, which
166 would require demand to be elastic at the optimum. We read it as a short-run object. Switching a
167 production workload to another model requires re-testing prompts, harnesses and evaluation suites,
168 so the relevant margin adjusts over months rather than days. It also rationalises the rigidity: a
169 provider that would gain little volume from a price cut, and lose little from holding steady, has weak
170 incentive to reprice a model already in the field.

171 4 Entry without displacement

172 The daily panel covers every model with usage for the four weeks to 20 August 2026, a window
173 containing 10 releases by providers with priced, scored models, of which 5 are open-weight. The
174 window is unusual in one respect that turns out to be the point: the market grew 65% over it.

175 The left panel of Figure 3 decomposes daily volume into incumbents and models released during
176 the window. Entrants reached 28% of platform tokens within four weeks, and 69% of that volume
177 went to open-weight entrants. Incumbent volume did not fall. It grew 18%. Entry accounts for 72%
178 of net growth, and the residual is incumbent expansion.

179 Aggregate volume can hide reallocation, so we test for it directly. For each release k at date T_k we
180 stack all incumbents j that were already serving traffic, and estimate

$$\log y_{jt} = \beta \mathbf{1}\{t \geq T_k\} \times \tau_{jk} + \delta_{jk} + \lambda_{kt} + X'_{jt}\theta + \varepsilon_{jkt}, \quad (2)$$

181 where δ_{jk} is an incumbent-by-entrant effect and λ_{kt} an entrant-by-day effect, so that β compares
182 incumbents exposed to the same release on the same day. The exposure measure $\tau_{jk} = S_{jk} r_{jk}$




Figure 3: Entry is absorbed by growth. Left: daily platform volume split between models already serving traffic at the start of the window and models released during it. Right: event-time coefficients from (2) on log incumbent volume, per unit of exposure to the entrant, with 95% intervals clustered by incumbent.

combines technological proximity with the entrant’s relative value for money: $S_{jk} = \exp(-\kappa d_{jk})$ with d_{jk} the Mahalanobis distance between capability vectors, and $r_{jk} = \log(q_k/p_k) - \log(q_j/p_j)$. Because capability and price both trend, we also control for post-release interactions with the incumbent’s own capability and price.

We find no diversion. The estimate of β is 0.01 (s.e. 0.07), and the event-time path in the right panel of Figure 3 is flat on both sides of the release. The confidence interval excludes differential declines larger than 0.13 log points per unit of exposure, and tightens to 0.10 and 0.04 under the Euclidean and cosine distance measures reported in Table 4. A design with ten releases over four weeks cannot detect small effects, and we do not claim that diversion is zero. Nor can this design rule out that exposure is measured with error: if capability similarity is a poor proxy for substitutability, the null is uninformative about diversion even though it is informative about the aggregate pattern. What we can say is that in a market growing at this rate, the displacement channel is small relative to the expansion channel, which is the comparison that matters for how the constraint operates.

A longer horizon. Four weeks is short, and displacement plausibly takes quarters. The archived panel lets us look further out. We take the 19 open-weight releases in it that reached at least one percent of platform tokens, and compare aggregate volume of models already in the market over the eight weekly observations before and after each. Incumbent volume fell in 32% of these episodes; the median change was +0.25 log points against a median market change of +0.52. The release of DeepSeek R1 is the sharpest case: platform volume rose 0.56 log points over the sixteen weeks around it while incumbent volume was flat at −0.03. Incumbents lose share, because the market grows around them, but for the most part they do not lose volume.

5 What open weights are worth

Because posted prices do not respond and incumbents are not displaced, the value of open-weight competition has to be measured on the choice set rather than on prices. We estimate a nested-logit demand system over the 172 models with prices and capability scores, nesting models by capability tercile so that models of similar capability are closer substitutes than models drawn from different tiers. Writing $g(j)$ for the nest of model j and $\sigma \in [0, 1)$ for the within-nest correlation [Cardell, 1997, Berry, 1994], shares take the form

$$s_j = \frac{\exp(\delta_j/(1-\sigma))}{D_{g(j)}} \cdot \frac{D_{g(j)}^{1-\sigma}}{1 + \sum_h D_h^{1-\sigma}}, \quad D_g = \sum_{k \in g} \exp(\delta_k/(1-\sigma)), \quad (3)$$

with mean utility $\delta_j = \beta_p \log p_j + \alpha' x_j + \xi_j$, where x_j collects the capability index, context length, modality, reasoning support, latency and age. The outside option is the fringe of models we cannot score, which holds 2.4% of tokens.

Two features of (3) make the exercise tractable. Observed shares invert to mean utilities, $\delta_j = \log s_j - \log s_0 - \sigma \log s_{j|g(j)}$ [Berry, 1994], so the system reproduces the data exactly and no part of the counterfactual depends on how well the characteristics fit. And the counterfactual depends only on (β_p, σ) : deleting the open-weight models from the choice set changes the denominators in (3), redistributing their volume across the remaining models and the outside option.

Choosing the parameters. We do not attempt to identify β_p and σ from cross-sectional share variation, because the instruments this literature relies on fail in this market. The differentiation instruments of Gandhi and Houde [2019] are functions of the distribution of rival characteristics, which here is a distribution over an essentially one-dimensional capability index, so they are near-collinear with a model’s own capability and therefore with its unobserved quality. The obvious cost shifter, parameter count, is available only for open-weight models and predicts unmeasured quality about as strongly as it predicts serving cost: in our data it correlates 0.65 with the capability index, and instrumenting with it moves the price coefficient from -0.79 under ordinary least squares to -0.27 despite a first-stage F above 900, which is the signature of a violated exclusion restriction rather than of a corrected bias. We therefore fix $\beta_p = -1$ and report every result across a range of σ . That value is deliberately conservative: it exceeds in magnitude both the within-model estimate of Section 3.3 and the adjustment-corrected long-run elasticity of about -0.6 implied by (4) below, and it is close to the value reported by Demirer et al. [2025]. Because surplus enters as a log-sum divided by $|\beta_p|$, a larger magnitude yields a smaller dollar figure, so the choice understates rather than inflates what open weights are worth.

Results. Removing open-weight models from the choice set at posted prices, the average price paid per token rises from \$0.70 to \$1.80, an increase of 156%. Consumer surplus, computed as the log-sum expected utility of Small and Rosen [1981], falls 21%, worth about \$2.4m per day at the platform’s current volume. The percentage change in surplus does not depend on β_p , which enters the baseline and the counterfactual identically; only the conversion into dollars scales with it. The loss is larger when substitution within a capability tier is weaker, ranging from \$3.9m to \$1.3m per day as σ moves from 0 to 0.7 (Figure 4 and Table 5). It also depends on how large the outside option is taken to be, which is a modelling choice rather than a measurement: raising the outside share from the observed 2.4% to 50% raises the proportional surplus loss, from 21% to 47%, while lowering the dollar figure to \$1.0m per day. The price effect is invariant to both, because it is governed by the composition of the surviving choice set rather than by the strength of substitution. Concentration moves against the presumption that open weights fragment the market: the provider Herfindahl index rises from 897 to 1048 when open models are removed, because open-weight volume is spread across more suppliers than the proprietary volume that replaces it.

These numbers hold prices fixed. That is the right convention given the rigidity documented in Section 3.2, and it makes the exercise a statement about variety rather than about pricing conduct, but it understates the long-run effect. Allowing proprietary providers to reprice would require a pricing model, and the elasticity in Section 3.3 is not consistent with the static Bertrand model that would ordinarily supply one, so we leave that margin unpriced and read our estimates as a lower bound.

The left panel of Figure 4 repeats the exercise separately in each of the 29 task-level markets, where shares of spending are reported directly. The implied price increase rises almost one for one with the open-weight share of task spending, with a slope of 0.98. Averaged across tasks and weighted by spending, prices rise 18% and surplus falls 13%, with the largest effects in agentic and coding work, where open-weight models hold about a third of spending, and the smallest in data extraction and transformation, where they hold under a sixth. The task-level numbers are smaller than the aggregate because task shares are measured on spending, which weights the expensive proprietary models that open weights do not displace.

6 Discussion

The three facts fit together. Short-run demand for a given model is inelastic, so providers gain little by cutting the price of a model in the field and do not; because posted prices do not move, competition is expressed in the terms on which new models enter; and because the market is growing quickly, entrants can take a large share of volume without taking it from anyone. In that environment,




Figure 4: The value of the open-weight choice set. Left: for each of the 29 task markets, the price increase implied by removing open-weight models against their share of task spending; marker area is proportional to the task’s share of platform spending. Right: aggregate surplus loss per day across values of the within-tier substitution parameter and the price coefficient.

the standard empirical signature of competitive discipline, a price response by incumbents, is absent even though the constraint is economically large. Our counterfactual puts the value of the open-weight choice set at roughly one fifth of consumer surplus on this platform.

A simple account. The three fit together through the share of demand that is actually in play. Let M_t be total demand, growing at rate g per period, and let a fraction ϕ of existing demand be revisited each period, whether because a workload is rebuilt or because a team re-evaluates its choice. Contestable demand is $(g + \phi)M_t$, and a model’s volume evolves as

$$y_{j,t+1} = (1 - \phi) y_{jt} + \omega_j (g + \phi) M_t, \quad (4)$$

where ω_j is its share of contestable demand. Three implications follow. An incumbent loses volume only when $\omega_j (g + \phi) M_t < \phi y_{jt}$, that is, when its share of contestable demand falls below $\phi / (g + \phi)$ times its share of the stock; with $g \approx 0.13$ per week in our window, an incumbent with an average share of new demand keeps growing. An entrant reaches a share of about $\omega_k (g + \phi) / (1 + g)$ after one period, so entrants can take a large share of volume quickly without anyone shrinking. And a permanent price change reaches only the contestable margin, so the response measured τ periods out is attenuated by roughly $1 - (1 - \phi - g)^\tau$ relative to its long-run value.

That last point bounds how much of the inelasticity is an artefact of slow adjustment. Taking $\phi + g \approx 0.23$ per week and the six weekly observations over which Figure 2 shows the response settling, the attenuation factor is about 0.8, so the long-run elasticity implied by our estimate is near -0.6 . Demand for a given model is inelastic even after adjustment. The static first-order condition of a provider maximising current profit is therefore not satisfied at the posted price, which leaves two readings: providers price against a margin we do not observe, such as capacity, hardware amortisation or the option value of usage; or the posted price is not chosen to maximise current profit at all. The rigidity in Section 3.2 favours the second.

This also reconciles two readings of the existing evidence. Nagle and Yue [2025] find that a great deal of proprietary usage sits on models dominated by an open alternative, and infer large foregone savings. That pattern is a puzzle if demand adjusts quickly; it is what one should expect if the short-run elasticity is about one-half and switching a workload is costly. The gains they identify are real, but they accrue as workloads turn over rather than as incumbents reprice.

Limitations. Our market is one router, and one that selects for price-sensitive developers, so the elasticities and the counterfactual describe the segment where substitution toward open weights is easiest. The capability index is built from third-party benchmarks that compress into a single dimension, and models that look identical on it are plainly not interchangeable in practice, which is precisely why the discrete-choice approach is needed and also why the index should not be read as a complete description of quality. We do not identify β_p and σ from the data and instead report

sensitivity across a range; a design with exogenous variation in the choice set, such as a large release observed over a longer window, would pin them down. The diversion test has power against moderate effects only. Finally, the counterfactual holds prices fixed, and is therefore silent about how the proprietary frontier would be priced in a world without open weights.

Implications. For policy, the result reframes what to measure. Debates about the competitive effect of open weights have proceeded on the presumption that the effect, if real, would show up in the prices proprietary providers charge. In a market where posted prices are set once and left alone, that presumption is wrong, and an assessment based on incumbent price responses will conclude that open weights do not matter. The relevant margin is the price at which the next unit of demand is served, and on that margin the constraint is substantial.

References

- Y. Bengio et al. International AI Safety Report 2026. arXiv:2602.21012, 2026.
- S. Berry. Estimating discrete-choice models of product differentiation. *RAND Journal of Economics*, 25(2):242–262, 1994.
- S. Berry, J. Levinsohn, and A. Pakes. Automobile prices in market equilibrium. *Econometrica*, 63(4):841–890, 1995.
- M. Bils and P. J. Klenow. Some evidence on the importance of sticky prices. *Journal of Political Economy*, 112(5):947–985, 2004.
- N. S. Cardell. Variance components structures for the extreme-value and logistic distributions with application to models of heterogeneity. *Econometric Theory*, 13(2):185–213, 1997.
- A. Cavallo. Scraped data and sticky prices. *Review of Economics and Statistics*, 100(1):105–119, 2018.
- M. Demirel, A. Fradkin, N. Tadelis, and S. Peng. The emerging market for intelligence: Pricing, supply, and demand for LLMs. NBER Working Paper 34608, 2025.
- A. Gandhi and J.-F. Houde. Measuring substitution patterns in differentiated-products industries. NBER Working Paper 26375, 2019.
- A. Korinek and J. Vipra. Concentrating intelligence: Scaling and market structure in artificial intelligence. *Economic Policy*, 40(121):225–256, 2025.
- F. Nagle and D. Yue. The latent role of open models in the AI economy. Working paper, 2025.
- E. Nakamura and J. Steinsson. Five facts about prices: A reevaluation of menu cost models. *Quarterly Journal of Economics*, 123(4):1415–1464, 2008.
- A. Nevo. Measuring market power in the ready-to-eat cereal industry. *Econometrica*, 69(2):307–342, 2001.
- K. A. Small and H. S. Rosen. Applied welfare economics with discrete choice models. *Econometrica*, 49(1):105–130, 1981.

A Data construction

Endpoints. The catalogue endpoint returns, for every listed model, the posted input and output price per token, context length, input and output modalities, a creation timestamp, and the Hugging Face repository identifier when the weights are public. The leaderboard endpoint returns tokens, requests, tool calls and media counts by model and serving variant. A benchmark endpoint returns third-party capability scores and usage-weighted effective prices. A task endpoint returns, for each of 29 usage categories, the category’s share of platform spending and the share held by each of its ten largest models, in both spending and token terms. An endpoint listing gives each hosting provider’s price, quantisation and uptime for a given model. All are public and unauthenticated.

Daily usage. Each model’s public page embeds its own trailing-month daily usage series in the page’s streamed render payload. We request one page per model with observed usage and parse the payload, which yields 457 models over 30 days. The same payload carries per-endpoint latency and throughput percentiles, which we aggregate to the model level by request count.

Historical coverage. Internet Archive captures supply history. Captures of the catalogue endpoint give a price and characteristic panel from July 2023. Captures of the leaderboard page embed the leaderboard payload in the same streamed format, so we recover model-level usage without relying on the rendered page. Two vintages are present and we treat them separately: captures before late 2024 embed a series of non-overlapping weekly buckets, and later captures embed a single trailing-window snapshot. Shares are comparable across vintages; levels are not, so all specifications include date fixed effects, and Table 3 reports the two vintages separately. We fetch captures at a fixed spacing with adaptive backoff, which is the binding cost of the collection: the archive requests take a few hours, and every other step runs in minutes on a laptop.

Merging. Usage and price observations are matched by model identifier and nearest date within ten days. Openness is taken from the current catalogue where the model is still listed, because the weight-repository field is absent from the earliest captures. Prices are dollars per million tokens throughout. The fixed-weight price index used in (1) weights posted input and output prices by the sample-mean token mix, so that it carries no feedback from a model’s own usage composition; Table 3 shows the estimate is not sensitive to that choice.

B Capability imputation

Benchmark scores are missing for smaller models. For each missing dimension we fit a ridge regression on the dimensions the model does have, using models scored on both, and predict the missing value. Masking half of the fully scored models and predicting the held-out values gives R^2 of 0.99, 0.97 and 0.96 for the reasoning, coding and agentic dimensions, with mean absolute errors of 1.7, 5.0 and 3.1 points against standard deviations of 16.4, 22.6 and 18.5. Imputation raises coverage from 109 models and 77% of tokens to 180 models and 91% of tokens. All results in the main text hold on the directly scored subsample.

C Robustness

Table 3 reports the elasticity across price measures, outcome measures, capture vintages, openness and sample period. The estimate stays between -0.38 and -0.63 . The realised blended price, which mixes posted prices with the model’s own input-output composition, gives the largest magnitude, as expected if a shift toward prompt-heavy workloads lowers the average price per token and raises volume at the same time; the fixed-weight index removes that channel and is our preferred measure.

Two further checks bear on the elasticity. A price cut may accompany the launch of a replacement by the same provider, which would depress the old model’s volume for reasons other than price. Adding an indicator for observations within 28 days of a same-provider release, which covers 38% of the panel, leaves the estimate at -0.534 ; dropping those observations gives -0.499 (s.e. 0.123).

Table 4 reports the diversion test under three definitions of technological distance. None of the six estimates is distinguishable from zero. The Euclidean and cosine measures give tighter standard errors than the Mahalanobis measure, and bound the differential decline for the most exposed incumbents at 0.13 and 0.09 log points respectively.

D Counterfactual sensitivity

Table 5 reports the counterfactual across the within-tier substitution parameter, and robustness. `json` reports it across the size of the outside option as well. The price effect is nearly invariant to σ , because it is driven by the composition of the surviving choice set rather than by the strength of substitution. The surplus loss falls as σ rises, because closer substitutes within a capability tier make the proprietary models better replacements for the open-weight models that leave.

Table 3: Own-price elasticity across specifications. Each row is a separate estimate of (1) with model and date fixed effects; standard errors clustered by model.

Specification	Estimate	Std. error	Observations	Models
Fixed-weight price (baseline)	−0.534	0.102	19,803	634
Input price	−0.526	0.096	19,803	634
Output price	−0.485	0.109	19,803	634
Realised blended price	−0.658	0.095	19,803	634
Requests rather than tokens	−0.463	0.087	19,803	634
Weekly captures only	−0.472	0.123	15,133	432
Trailing captures only	−0.476	0.125	4,670	428
Open-weight models	−0.582	0.116	8,873	321
Proprietary models	−0.534	0.204	10,930	316
2025 onwards	−0.527	0.100	7,906	542

Table 4: Diversion at entry under alternative distance measures. Each row estimates (2) with incumbent-by-entrant and entrant-by-day fixed effects and post-release controls for the incumbent’s own capability and price.

Distance	Exposure	Estimate	Std. error	Observations
Mahalanobis	similarity $\times$ relative value	0.010	0.070	11,820
Mahalanobis	similarity	0.019	0.120	11,820
Euclidean	similarity $\times$ relative value	−0.006	0.047	11,820
Euclidean	similarity	0.008	0.065	11,820
Cosine	similarity $\times$ relative value	0.022	0.031	11,820
Cosine	similarity	0.083	0.086	11,820

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and Section 1 state the three measurement results and the counterfactual, each of which is established in Sections 3–5, and both state that the market studied is one inference router.

Guidelines:

- The answer [\[N/A\]](#) means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [\[No\]](#) or [\[N/A\]](#) answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: Section 6 gives a limitations paragraph covering market coverage, the capability index, the calibrated demand parameters and the power of the diversion test.

Guidelines:

- The answer [\[N/A\]](#) means that the paper has no limitation while the answer [\[No\]](#) means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.

Table 5: The open-weight counterfactual across substitution parameters, at $\beta_p = -1$. Prices are dollars per million tokens; the surplus loss is evaluated at the platform’s volume on the reference day.

σ	Price paid (\$/M)	Price without open (\$/M)	Surplus change (%)	Surplus loss (\$m per day)	HHI (without open)
0.0	0.70	1.82	−33.8	3.9	1109
0.2	0.70	1.81	−27.5	3.2	1082
0.4	0.70	1.80	−21.0	2.4	1048
0.6	0.70	1.79	−14.3	1.7	1006
0.7	0.70	1.78	−10.8	1.3	982

- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper contains no theorems. The assumptions of the demand model are stated in Section 5.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Every input is a public endpoint or an Internet Archive capture; Appendix A documents the collection and construction, and the released code reproduces the panels and all reported numbers.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The collection and analysis code is released, and the raw inputs can be rebuilt from the public endpoints and archive captures the code queries.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).

- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Sample construction, variable definitions, fixed effects and clustering are given in Sections 3–5 and Appendix A.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: All regressions report standard errors clustered by model, and the calibrated demand parameters are reported across a grid rather than at a single value.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: All estimation runs in a few minutes on a laptop; the only substantial cost is the rate-limited archive collection, which is documented in Appendix A.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research conforms with the NeurIPS Code of Ethics. It uses aggregate, model-level public data and no personal data.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Section 6 discusses the implication for how the competitive effect of open weights should be assessed.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.

- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper releases analysis code and no models, weights or scraped personal data.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The data sources are public API and web endpoints and Internet Archive captures; Appendix A records the sources and access terms.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The released code is documented, and Appendix A describes the derived panels it produces.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The paper does not involve crowdsourcing or human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The paper does not involve human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [Yes]

Justification: Large language models were used for coding assistance and for editing prose. They were not part of the research method: the data collection, the estimation and the reported results were specified and checked by the authors.

Guidelines:

724
725
726
727

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.